
Error Structure Determines Correctness Preference in Contradictory Training Data

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 When transformers train on contradictory data—the same problem with both cor-
2 rect and incorrect solutions—which answer do they prefer? We hypothesize that
3 next-token prediction, as a compression process, favors whichever answer cluster
4 has lower description length; truth benefits only when errors lack internal struc-
5 ture. We test this by training transformers (3.5M–1B parameters) from scratch
6 on controlled corpora, systematically varying the structure of errors. We find that
7 (a) when errors are random, models develop a correctness preference scaling from
8 65% to 85% with model size; (b) when errors follow a single coherent alternative
9 rule, this preference vanishes ($\sim 45\text{--}51\%$); (c) two competing wrong rules suffice
10 to restore it ($47\% \rightarrow 78\%$). The pattern reproduces on Wikipedia paragraphs
11 with entity substitution (71% vs 46%) and at 1B scale on a mixed natural-text
12 corpus (77% vs 47%). These results are consistent with the hypothesis that, in
13 controlled contradictory corpora, model preference tracks the relative compress-
14 ibility of competing answer systems rather than truth per se.

15 1 Introduction

16 Real-world training corpora are noisy: the same question may receive contradictory answers across
17 different documents. Yet language models trained on such data tend to prefer correct information.
18 If we do not understand *why* this filtering works, we cannot predict *when* it will fail—for instance,
19 when errors are not diverse but internally consistent.

20 Several explanations have been offered: scaling improves factual performance [Kadavath et al.,
21 2022]; frequency matters, since factual accuracy correlates with how often correct information
22 appears in training data [Elazar et al., 2022, Joshi et al., 2024, Kandpal et al., 2023]; and truth-
23 correlated internal representations have been found [Burns et al., 2023, Marks and Tegmark, 2024,
24 Bürger et al., 2024]. Kalai and Vempala [2024] established that calibrated models must hallucinate
25 at a rate bounded by the corpus’s monofact rate. Yet none of these explain why the *training objective*
26 *alone* would favor one answer over another when both appear at equal frequency in the same format.

27 We propose that the answer lies in the structure of errors. Minimizing cross-entropy is equivalent to
28 minimizing code length [Shannon, 1948, Delétang et al., 2024], linking LLM training to the Mini-
29 mum Description Length principle [Rissanen, 1978, Grünwald, 2007]. In our task families, a correct
30 rule system is compact; diverse random errors must be memorized individually. But a *coherent* false
31 system—internally consistent, just wrong—can be equally compact, and the preference vanishes.

32 We test this hypothesis in a controlled experimental setting. We train transformers (3.5M–1B pa-
33 rameters) from scratch on corpora where each mathematical problem appears with both correct and
34 incorrect solutions, systematically varying the structure of errors. This design isolates error structure
35 as the independent variable while holding frequency, format, and domain constant.

Our contribution is a single experimentally supported hypothesis: **in controlled contradictory corpora, the compression objective favors consistency, not truth.** We show that (a) random errors produce a correctness preference scaling from 65% to 85%; (b) a single coherent alternative rule eliminates this preference; (c) two competing rules restore it, pinpointing the compressibility boundary; (d) the same contrast reproduces on Wikipedia paragraphs with entity substitution; and (e) the effect persists at 1B parameters on a mixed natural-text corpus (Section 4.7). Whether this extends to large-scale pretraining remains open (Section 5).

2 Related Work

Prediction as compression. The link between prediction and compression traces to Shannon [1948] and was formalized by Solomonoff [1964] and Rissanen’s MDL principle [Rissanen, 1978, Grünwald, 2007]. Delétang et al. [2024] showed that LLMs are universal compressors. Huang et al. [2024] discovered a linear correlation ($r \approx -0.93$) between compression quality and benchmark performance. Wan and Mei [2025] proved that LLM training approximates Solomonoff induction. Pan et al. [2025] linked compression to knowledge acquisition and scaling; Chlon et al. [2025] linked compression failures to hallucination. Contemporaneous work by Marty et al. [2026] formalizes simplicity bias through MDL, and Aksenov et al. [2026] connects compression to mathematical modeling. We build on the MDL framework by experimentally varying the description length of false answer systems.

Internal representations and world models. Compression can give rise to structured internal representations: Li et al. [2023a] found board representations in an Othello model, Gurnee and Tegmark [2024] discovered linear space-time representations in Llama-2, and several studies found truth-correlated representations [Marks and Tegmark, 2024, Burns et al., 2023, Li et al., 2023b, Azaria and Mitchell, 2023, Ravfogel et al., 2025]. Our work operates at the behavioral level: we identify data-level conditions under which compression produces a preference for correct completions, leaving activation-level analysis for future work.

Truthfulness and data statistics. Joshi et al. [2024] linked truthfulness to “personas” in pretraining data. Elazar et al. [2022] demonstrated frequency dependence. Kandpal et al. [2023] showed a direct relationship between document count and accuracy. Li et al. [2024] studied LLM preferences under conflicting knowledge and found that formality and surface quality predict which version the model favors—a consistency-driven explanation complementary to ours. A growing body of work studies knowledge conflicts directly: Xie et al. [2024] surveyed approaches to resolving conflicts between parametric and contextual knowledge, and Longpre et al. [2021] showed that entity-substituted evidence can override parametric memory. Unlike these analyses, we fix frequency, format, and surface quality, varying only the compressibility of the error system itself—isolating error structure as the causal variable rather than studying post-hoc conflict resolution.

Simplicity bias and noisy labels. Neural networks prefer simple functions [Valle-Perez et al., 2019, Mingard et al., 2021, Goldblum et al., 2024]. The noisy labels literature directly parallels our setup: Zhang et al. [2017] showed memorization of random labels; Rolnick et al. [2017] showed robustness to massive label noise. Grokking connects to compression as a memorization-to-generalization transition [Nanda et al., 2023, Liu et al., 2023]. Our experiments extend this line by showing that “structured noise”—coherent errors—is not filtered out. To our knowledge, systematic variation of error compressibility in a denoising setting has not been studied before.

3 Methodology

Overview. We test the hypothesis that next-token prediction favors whichever answer system has lower description length, regardless of its truth value. To isolate this, we train models from scratch on controlled corpora where each problem appears with contradictory answers, varying only the structure of errors. The core experiments compare random errors (high description length) against coherent errors (low description length) at equal frequency (Section 4.1), then probe the boundary with multi-rule errors of intermediate compressibility (Section 4.2). Transfer to natural language is tested via Wikipedia entity substitution (Section 4.3).

3.1 Hypothesis and MDL Framing

Consider a corpus where each problem appears with a correct answer (theory T_1) and an alternative (T_2). An MDL learner [Rissanen, 1978, Grünwald, 2007] minimizes $L(M) + L(D|M)$. When $K(T_2) \gg K(T_1)$ (random errors), the false system’s description length grows with corpus size, favoring T_1 . When $K(T_2) \approx K(T_1)$ (coherent errors), both are compact and the learner has no basis to prefer one over the other. With N competing false rules, a random “selector” (which rule applies to which problem) adds $\log N$ bits per problem, progressively degrading the false cluster’s compressibility.

Throughout, “truth” means correctness of mathematical derivations and factual accuracy of Wikipedia paragraphs—models compress text, not reality. The compressibility gap is domain-dependent and smaller in natural language than in formal math (Section 4.3). Additional limitations are discussed in the Limitations section.

3.2 Models and Training

Custom decoder-only transformers with GPT-2-like architecture (pre-norm, GELU activation, causal mask). Size labels below are internal designations and do not correspond to OpenAI’s GPT-2 checkpoints (e.g., our “large” at 86M is smaller than GPT-2 Small at 117M):

Config	Layers	d_{model}	Heads	Parameters
tiny	4	256	4	~3.5M
small	6	384	6	~12M
medium	8	512	8	~26M
large	12	768	12	~86M

Table 1: Model configurations used in all experiments.

Optimizer: AdamW (weight_decay=0.01), cosine decay with linear warmup, lr = $3\text{e-}4$, seq_len = 256, batch_size = 32, 5000 steps. All experiments use 4 random seeds for core conditions (2 where noted). Learning curves confirm behavioral plateau by step 3000–4000 at all sizes (Appendix F). Denoising experiments use PyTorch; multi-rule and Wikipedia experiments use MLX. All problems are generated and verified by SymPy; no human annotation is involved. Full generation templates are provided in the supplementary material.

3.3 Corpus Design

Denoising setup (primary). A generator creates mathematical problems of four types: multi-step arithmetic, factoring, equation solving, and differentiation, formatted as step-by-step derivations in English. The key design: **each problem appears with contradictory answers**—modeling the scenario where the same question receives conflicting responses. The tokenizer is character-level (vocab ~57); BPE robustness is verified in Section 4.4.

Condition	Correct	Incorrect	Ratio	Purpose
Equal random	1	1 random	1:1	Signal extraction from random noise
Equal coherent	1	1 coherent	1:1	Control: coherent errors eliminate bias
2:1 noise	1	2 random	1:2	Noise tolerance at moderate noise
4:1 noise	1	4 random	1:4	Noise tolerance at high noise

Table 2: Denoising corpus conditions. Each corpus contains 5,000 unique problems.

Wikipedia setup (transfer). To test generalization beyond formal math, we construct corpora from 20,000 Wikipedia paragraphs. Using spaCy NER, we create two corruption modes: *random substitution* (each entity replaced with a random entity of the same type) and *coherent substitution* (a consistent global mapping, e.g., every “France” $\rightarrow$ “Japan”, every “Paris” $\rightarrow$ “Tokyo”). Coherent substitution preserves grammatical structure and cross-entity consistency within each paragraph;

Size	Params	Random Accuracy	Coherent Accuracy	Random Seeds	Coherent Seeds
tiny	3.5M	65.3% $\pm$ 1.3%	43.5% $\pm$ 2.6%	4	4
small	12M	74.6% $\pm$ 1.6%	44.5% $\pm$ 3.0%	4	4
medium	26M	81.1% $\pm$ 1.2%	45.8% $\pm$ 3.4%	4	4
large	86M	85.2% $\pm$ 2.3%	51.0% $\pm$ 0.8%	2	2

Table 3: Denoising paired evaluation: equal random vs equal coherent errors, across model sizes. All $\pm$ values are standard deviations across random seeds. Random accuracy scales with model size; coherent accuracy remains near chance.




Figure 1: The central contrast. Random errors: accuracy scales from 65% to 85%. Coherent errors: accuracy stays near chance across all sizes.

119 random substitution breaks co-occurrence patterns (e.g., a paragraph may mention “Kumamoto”
120 in a context about French history). Both modes substitute entities of matching type, so surface-
121 level grammatical anomalies (gender, case agreement) are minimal. Models train on 50/50 mixes;
122 evaluation uses 2,000 held-out paired paragraphs.

123 3.4 Evaluation

124 **Paired evaluation (primary).** For each problem, a shared prompt is generated with two com-
125 pletions (correct and incorrect). NLL is computed on completion tokens only, conditioned on the
126 prompt. The primary metric is **pair accuracy**: the fraction of pairs where the model assigns lower
127 NLL to the correct completion. This is equivalent to the Common Language Effect Size (CLES;
128 McGraw and Wong, 1992). Length-matched evaluation confirms that accuracy is not driven by
129 completion length differences (within 1 pp across all conditions).

130 **Generative evaluation.** Greedy decoding on 500 test prompts with automated SymPy verification.
131 This confirms that the discriminative effect extends to generation (Section 4.5; Table 6).

132 4 Results

133 4.1 The Central Contrast: Random vs Coherent Errors

134 When the same problem appears with both a correct and a random wrong answer, the model learns
135 to prefer the correct one—accuracy reaches 85% at 86M parameters (Table 3, Figure 1). When the
136 wrong answer follows a coherent rule system, the effect disappears: accuracy hovers near 50% at all
137 scales.

138 Random accuracy increases monotonically with scale (65% $\rightarrow$ 75% $\rightarrow$ 81% $\rightarrow$ 85%). Coherent
139 accuracy remains near chance across all scales, consistent with the MDL prediction that equally
140 compressible systems at equal frequency should be indistinguishable. The below-chance coherent
141 accuracy at tiny ($\sim$ 43%) is consistent with the textual simplicity of the false rules (e.g., dropping
142 terms in derivatives); the compressor may prefer whichever system is simpler to encode. The multi-
143 rule experiment (Section 4.2) controls for this surface-form effect.

N Rules	Accuracy	p
1 (coherent)	46.6% $\pm$ 2.4%	~ 1.0
2	77.6% $\pm$ 1.3%	$< 10^{-6}$
3	82.8% $\pm$ 0.1%	$< 10^{-6}$
5	84.8% $\pm$ 0.5%	$< 10^{-6}$
10	88.3% $\pm$ 0.7%	$< 10^{-6}$

Table 4: Multi-rule paired evaluation (tiny, 3.5M, standard 50/50, 4 seeds). The $N=1$ baseline (46.6%) differs from the denoising coherent accuracy in Table 3 (43.5%) because multi-rule experiments use the standard (non-denoising) setup. $N=2$ on small (12M) yields 86.3% $\pm$ 0.8%, confirming the effect strengthens with capacity.




Figure 2: Multi-rule crossover. Accuracy jumps from 47% ($N=1$) to 78% ($N=2$), then grows gradually through $N=10$.

144 4.2 Multi-Rule Errors: The Sharp Crossover

145 The denoising experiments establish two poles: one coherent rule yields chance, random errors
 146 yield strong bias. We probe the boundary by training on corpora with N alternative wrong rules per
 147 task type, where each problem’s rule is chosen at random. Each rule is compact, but the mapping
 148 “problem $\rightarrow$ rule” is unpredictable.

149 This is the most mechanistically revealing result (Table 4, Figure 2). With one false rule, $K(T_2) \approx$
 150 $K(T_1)$. With two rules, the learner must encode a random selector—which rule applies to which
 151 problem—adding high-complexity bits that break compressibility. At $N=10$, accuracy (88.3%)
 152 exceeds even the random baseline (79.5%), because multi-rule errors combine compactness of indi-
 153 vidual rules with incompressibility of the selector.

154 4.3 Transfer to Wikipedia Text with Entity Substitution

155 The random/coherent contrast reproduces (Table 5): random substitution yields 70–71% accuracy
 156 (all $p < 10^{-6}$), coherent substitution yields 46–49% (at or below chance). The effect is weaker
 157 than in math (71% vs 85%) because natural language provides more textual flexibility for locally
 158 fluent errors. Unlike math, Wikipedia accuracy saturates near 70% across the tested range—the
 159 compressibility gap in natural language does not substantially benefit from additional capacity at
 160 this scale. Per-entity-type accuracy ranges from 82% (LOC) to 61% (CARDINAL), with geographic
 161 entities showing the strongest effect (Appendix E).

Size	Random Accuracy	Coherent Accuracy
tiny (3.5M)	69.6% $\pm$ 0.1%	48.7% $\pm$ 0.5%
small (12M)	70.7% $\pm$ 0.4%	46.6% $\pm$ 0.4%
medium (26M)	71.5% $\pm$ 0.8%	46.4% $\pm$ 0.8%
large (86M)	71.4% $\pm$ 0.7%	45.9% $\pm$ 1.5%

Table 5: Wikipedia entity substitution (50/50, 4 seeds per size). The random/coherent contrast reproduces on real text.

Size	Random-trained	Coherent-trained	Gap
tiny (3.5M)	30.5% $\pm$ 1.7%	20.8% $\pm$ 3.6%	+9.7 pp
small (12M)	46.7% $\pm$ 1.9%	31.7% $\pm$ 1.8%	+15.0 pp
medium (26M)	50.5% $\pm$ 1.3%	36.1% $\pm$ 2.7%	+14.4 pp
large (86M)	52.8% $\pm$ 1.1%	35.6% $\pm$ 1.8%	+17.2 pp

Table 6: Generative accuracy (4 seeds, 500 problems each). The gap widens with model size.

162 4.4 Robustness Checks

163 **Tokenization.** The primary experiments use a character-level tokenizer (vocab $\sim$ 57). To verify
164 that the effect is not an artifact of tokenization, we repeat the core conditions with a BPE tokenizer
165 (SentencePiece, vocab=1000). The effect survives and strengthens: random accuracy increases from
166 65.3% to 75.9% in denoising (79.5% to 85.6% in standard); coherent accuracy remains at chance
167 (49.3% and 45.9% respectively). Full results in Appendix F.

168 **Standard (non-denoising) setup.** In the standard setup, each problem appears once with either
169 a correct or incorrect solution (no within-problem contradiction). Paired accuracy is higher (79.5–
170 88.3% vs 65.3–85.2%), but the gap closes with scale (from 14 pp at tiny to 3 pp at large). Critically,
171 the random/coherent contrast holds identically: standard coherent accuracy is 47–52%, at chance.
172 Full comparison in Appendix D.

173 **Frequency vs compressibility.** For random errors, truth bias persists even when correct examples
174 are a small minority: 67% paired accuracy at 10/90 correct-to-incorrect ratio (Appendix D). For
175 coherent errors, the pattern reverses—frequency dominates preference far more strongly: at 40/60
176 coherent, the model follows the majority with 72% preference for the false system; at 20/80, this
177 rises to 91%. This asymmetry is a direct prediction of the MDL framework: when both systems
178 compress equally well, frequency is the only tiebreaker.

179 **Matched-control ablation.** A potential confound is that coherent errors affect more derivation
180 steps per problem ($\sim$ 3 on average) than standard random errors (exactly 1). To control for this, we
181 generate “matched random” errors that corrupt the same number of steps as coherent errors but at
182 random positions with random values. Matched-random accuracy is 74.5% $\pm$ 0.4% (4 seeds)—lower
183 than standard random (79.5%, which corrupts only 1 step) but far above coherent (47.2%, which also
184 corrupts $\sim$ 3 steps). The 27 pp gap between matched-random and coherent at identical step counts
185 confirms that compressibility, not surface-level corruption load, drives the random/coherent contrast.

186 **Noise tolerance.** The correct signal persists under increasing noise ratios (1:2 and 1:4 random),
187 with accuracy degrading gracefully from 85% to 66–75% at large scale. Details in Appendix D.

188 4.5 Generative Evaluation

189 Paired evaluation is a forced-choice (discriminative) setting. To verify that the effect extends to
190 generation, we run greedy decoding with SymPy verification on 500 problems per model across all
191 sizes.

192 The gap widens with model size (10 pp at tiny $\rightarrow$ 17 pp at large), confirming that the discriminative
193 preference translates into a generative advantage (Table 6). The absolute generation accuracy is
194 substantially lower than paired accuracy (53% vs 85% at large). This gap reflects the fundamental
195 difference between discrimination and generation: the model may assign higher likelihood to correct

Condition	seed43	seed44	seed45	Mean
Random	85.3%	85.7%	89.3%	86.8%
Coherent	49.3%	52.8%	49.8%	50.6%

Table 7: Qwen3-0.6B paired evaluation (50/50, 3 seeds). The random/coherent contrast reproduces across architectures.




Figure 3: Full scaling results. GPT-2 family (3.5M–86M, math-only), Qwen3-0.6B (420M, different architecture, math-only), and a Qwen3-architecture model (~1B, FineWeb-Edu + 8% math).

196 completions while lacking the capacity to produce them reliably from scratch. Both metrics agree
197 on the direction: random-trained models outperform coherent-trained ones at all scales.

198 Per-task breakdown reveals that generation difficulty varies sharply by type. For random-trained
199 large models: derivative (79.5%), algebra (70.4%), equation (61.0%), arithmetic (0%). Multi-step
200 arithmetic requires exact carry propagation across many digits—a well-known difficulty for au-
201 toregressive generation at this scale—even though paired evaluation shows strong discrimination.
202 Coherent-trained models show lower generation accuracy across all types, with the gap largest for
203 algebra (70.4% vs 38.3%) and smallest for derivative (79.5% vs 56.5%).

204 4.6 Architecture Robustness: Qwen3

205 To verify that the effect is not specific to the GPT-2 architecture, we train a Qwen3-0.6B model
206 (420M non-embedding parameters, 28 layers, RoPE + GQA + SwiGLU + RMSNorm) from scratch
207 on the same denoising corpora. This architecture differs from GPT-2 in positional encoding,
208 attention mechanism, activation function, and normalization—testing whether the compression-
209 consistency effect is a general property of autoregressive transformers.

210 The pattern reproduces (Table 7, Figure 3): random errors yield 86.8% accuracy, coherent errors
211 yield 50.6% (chance). The random/coherent contrast is not specific to the GPT-2-like architecture
212 used in our primary experiments. Qwen3’s random accuracy (86.8%) is comparable to GPT-2 large
213 (85.2%) despite training on the same corpus sized for smaller models, suggesting further gains with
214 compute-matched training.

215 4.7 Scaling to 1B on Real Text

216 The preceding experiments use synthetic math corpora and models up to 86M parameters (420M for
217 architecture robustness). A natural objection is that the effect may be specific to toy-scale models or
218 synthetic data. To address both concerns simultaneously, we train a Qwen3-architecture model (~1B
219 parameters) from scratch on 1 GB of FineWeb-Edu [Penedo et al., 2024]—a curated subset of real
220 web text—with mathematical content filtered out and replaced by 8% contradictory math problems
221 (denoising format, equal random or coherent). This setup tests whether the compression-consistency

Condition	Accuracy	ΔLoss	p	Seeds
Random	76.8%	+0.098	$< 10^{-6}$	1
Coherent	46.7%	-0.003	~ 1.0	1

Table 8: Qwen3-architecture $\sim 1\text{B}$ trained on FineWeb-Edu + 8% contradictory math (denoising, 50/50). The random/coherent contrast reproduces at 1B scale on real text.

effect survives in a regime where (a) the model has two orders of magnitude more parameters than our primary experiments, and (b) the contradictory signal is a small fraction of a large, naturalistic corpus.

The result (Table 8) confirms that the random/coherent contrast holds: random errors yield 76.8% paired accuracy ($\Delta\text{Loss} = +0.098$, $p < 10^{-6}$), while coherent errors yield 46.7% (chance). The random accuracy is lower than the math-only GPT-2 large (85.2%) and Qwen3-0.6B (86.8%), likely because the contradictory math signal constitutes only 8% of the training corpus, diluted by general text. The key finding is that even in this diluted, naturalistic setting, the compression objective still distinguishes random from coherent errors—and the ΔLoss magnitude (+0.098) is the largest we observe, suggesting that the 1B model develops a stronger per-pair preference despite lower aggregate accuracy.

5 Discussion

Summary of findings. The experiments support a unified hypothesis: **in our controlled settings, the compression objective tracks consistency rather than truth.** In our task families, an internally consistent false rule system compresses comparably to the true one. Truth bias emerges only when false alternatives are structurally incoherent. The evidence converges across setups (denoising, standard, Wikipedia), error types (random, coherent, multi-rule), and evaluation modes (discriminative and generative).

The role of frequency. Our results reveal an asymmetry between random and coherent error regimes. For random errors, compressibility overrides frequency: truth bias persists even at 10/90 correct-to-incorrect ratio (67% accuracy). For coherent errors, frequency dominates: at 40/60 coherent mixing, the model follows the majority with 72% preference for the false system, rising to 91% at 20/80. This is consistent with the MDL prediction that when competing systems compress equally well, the more frequent one wins. The practical implication is that coherent falsehood needs no frequency advantage to neutralize truth bias; equal frequency suffices.

What the multi-rule experiment shows. The $N=1 \rightarrow 2$ transition provides strong evidence that compressibility, not surface-form complexity, is the operative variable. Each individual rule at $N=2$ remains compact, yet the random assignment of rules to problems introduces incompressible bits. This is difficult to reconcile with explanations based solely on lexical diversity or number of corrupted tokens.

Implications and scope. In our controlled settings, the training objective does not provide an inherent “truth compass.” Internally coherent falsehood remains competitive. Embedding cross-domain verification can partially restore truth bias, but the effect exhibits inverse scaling: verification accuracy drops from 71% at 3.5M to 61% at 86M parameters (Appendix F), as larger models absorb coherent within-domain patterns more readily. This inverse scaling is verified across 10 models (4 tiny + 4 small + 2 large) with consistent results. This has potential relevance for alignment (systematic false beliefs may resist filtering and grow harder to verify at scale), hallucinations (coherent misconceptions may persist; cf. Kalai and Vempala, 2024), and data curation (diverse errors are filtered effectively, but coordinated errors may not be). The 1B experiment on real web text (Section 4.7) provides initial evidence that the effect persists beyond toy scale and synthetic data, but all our experiments remain far below the scale of production LLMs. Whether these findings extend to large-scale pretraining on heterogeneous real-world corpora remains an open question. We view our results as identifying a mechanism and its boundary conditions, not as direct claims about production LLMs.

266 The pattern admits an analogy with Popper’s falsifiability criterion [Popper, 1959]: a false theory
267 requiring ad hoc corrections to fit observations has higher description length than a parsimonious true
268 one. The analogy is limited—the model does not test theories against evidence—but the structural
269 parallel between MDL-driven preference and falsifiability-driven theory selection is suggestive.

270 6 Conclusion

271 Whether the compression mechanism that produces correctness preferences in our experiments ex-
272 tends to large-scale pretraining remains an open question. In controlled experiments with transform-
273 ers up to 86M parameters—and up to $\sim 1\text{B}$ parameters on a mixed natural-text corpus (Section 4.7)—
274 we find that models prefer the correct answer only when errors are structurally incoherent; a single
275 coherent alternative rule eliminates this preference, and two competing rules restore it. The same
276 pattern reproduces on Wikipedia paragraphs with entity substitution, in generative evaluation, and
277 across architectures. In the settings we study, correctness preference is a compression artifact that
278 tracks the description length of the false answer system.

279 Limitations

280 **Model scale.** The primary experiments use models of 3.5M–86M parameters. Section 4.7 extends
281 to $\sim 1\text{B}$ parameters on a mixed natural-text corpus with a single seed per condition, providing a scale-
282 up sanity check but not a definitive large-scale confirmation. This remains far below production
283 LLMs (7B–400B+). Replication at multi-billion scale with compute-matched training and multiple
284 seeds remains an important next step.

285 **Domain specificity.** Mathematics provides an unusually crisp correct/incorrect distinction. The
286 effect weakens in natural language (71% vs 85%), where errors can remain locally fluent. Extending
287 to domains with competing real-world knowledge systems (e.g., conflicting news sources, historical
288 revisionism) remains open.

289 **Discriminative vs generative gap.** Paired accuracy (85%) substantially exceeds generative accu-
290 racy (53%) at large scale. This gap matters: real-world LLM behavior depends on generation, not
291 discrimination. Several factors may contribute: generation requires producing correct tokens autore-
292 gressively (errors compound), while paired evaluation only requires relative likelihood ranking. The
293 gap narrows with model size (49 pp at tiny, 32 pp at large), suggesting that scale partially bridges
294 it. Critically, both metrics agree on the direction and on the random/coherent contrast at all tested
295 scales—the discriminative preference is not an evaluation artifact. Still, paired accuracy should be
296 interpreted as a controlled diagnostic of model preference, not as a direct measure of output truthful-
297 ness.

298 **Seed counts.** Core conditions use 4 seeds; some conditions use 2. Sufficient for directional stabil-
299 ity, not for tight confidence intervals.

300 **Causal identification.** While we vary error structure as the independent variable, random and co-
301 herent corruptions may differ in additional ways (number of affected derivation steps, surface-form
302 simplicity). The multi-rule experiment partially addresses this by holding individual rule complexity
303 constant, but a fully matched control (identical corruption load, lexical diversity, and output length)
304 would strengthen the causal claim.

305 **Future directions.** Linear probing for “truth directions” vs “coherence directions” in models
306 trained under our conditions could connect behavioral findings to the internal representation liter-
307 ature. Testing interactions with RLHF and investigating whether the compressibility mechanism
308 interacts with in-context learning are natural extensions.

309 Ethical Considerations

310 This work demonstrates that, in controlled settings, internally consistent misinformation may be
311 harder for language models to filter than diverse errors. The finding does not directly demonstrate
312 vulnerability in production models. We believe transparent reporting of conditions under which truth
313 bias fails is more beneficial than concealment, as it can inform defensive measures in data curation
314 and model evaluation.

References

- V. Aksenov, E. Bodnia, M. H. Freedman, and M. Mulligan. Compression is all you need: Modeling mathematics. *arXiv preprint arXiv:2603.20396*, 2026.
- A. Azaria and T. Mitchell. The internal state of an LLM knows when it’s lying. In *Findings of EMNLP*, 2023.
- L. Bürger, F. A. Hamprecht, and B. Nadler. Truth is universal: Robust detection of lies in LLMs. In *NeurIPS*, 2024.
- C. Burns, H. Ye, D. Klein, and J. Steinhardt. Discovering latent knowledge in language models without supervision. In *ICLR*, 2023.
- L. Chlon, A. Karim, M. Chlon, and M. Awada. Predictable compression failures: Order sensitivity and information budgeting for evidence-grounded binary adjudication. *arXiv preprint arXiv:2509.11208*, 2025.
- G. Delétang, A. Ruoss, P.-A. Duquenne, E. Catt, T. Genewein, C. Mattern, J. Grau-Moya, L. K. Wenliang, M. Aitchison, L. Orseau, M. Hutter, and J. Veness. Language modeling is compression. In *ICLR*, 2024.
- Y. Elazar, N. Kassner, S. Ravfogel, A. Feder, A. Ravichander, M. Mosbach, Y. Belinkov, H. Schütze, and Y. Goldberg. Measuring causal effects of data statistics on language model’s factual predictions. *arXiv preprint arXiv:2207.14251*, 2022.
- M. Goldblum, M. Finzi, K. Rowan, and A. G. Wilson. Position: The no free lunch theorem, Kolmogorov complexity, and the role of inductive biases in machine learning. In *ICML*, 2024.
- P. D. Grünwald. *The Minimum Description Length Principle*. MIT Press, 2007.
- W. Gurnee and M. Tegmark. Language models represent space and time. In *ICLR*, 2024.
- Y. Huang, J. Zhang, Z. Shan, and J. He. Compression represents intelligence linearly. In *COLM*, 2024.
- N. Joshi, J. Rando, A. Saparov, N. Kim, and H. He. Personas as a way to model truthfulness in language models. In *EMNLP*, 2024.
- S. Kadavath, T. Conerly, A. Askell, T. Henighan, D. Drain, E. Perez, J. Kaplan, et al. Language models (mostly) know what they know. *arXiv preprint arXiv:2207.05221*, 2022.
- A. T. Kalai and S. S. Vempala. Calibrated language models must hallucinate. In *STOC*, 2024.
- N. Kandpal, H. Deng, A. Roberts, E. Wallace, and C. Raffel. Large language models struggle to learn long-tail knowledge. In *ICML*, 2023.
- J. Li, Y. Cao, S. Huang, and J. Chen. Formality is favored: Unraveling the learning preferences of large language models on data with conflicting knowledge. In *EMNLP*, 2024.
- K. Li, A. K. Hopkins, D. Bau, F. Viégas, H. Pfister, and M. Wattenberg. Emergent world representations: Exploring a sequence model trained on a synthetic task. In *ICLR*, 2023a.
- K. Li, O. Patel, F. Viégas, H. Pfister, and M. Wattenberg. Inference-time intervention: Eliciting truthful answers from a language model. In *NeurIPS*, 2023b.
- Z. Liu, Z. Zhong, and M. Tegmark. Grokking as compression: A nonlinear complexity perspective. *arXiv preprint arXiv:2310.05918*, 2023.
- S. Longpre, K. Perisetla, A. Chen, N. Ramesh, C. DuBois, and S. Singh. Entity-based knowledge conflicts in question answering. In *EMNLP*, 2021.
- S. Marks and M. Tegmark. The geometry of truth: Emergent linear structure in large language model representations of true/false datasets. In *COLM*, 2024.

358 T. Marty, E. Elmoznino, L. Gagnon, T. Kasetty, M. Nishikawa-Toomey, S. Mittal, G. Lajoie, and
359 D. Sridhar. A compression perspective on simplicity bias. *arXiv preprint arXiv:2603.25839*,
360 2026.

361 K. O. McGraw and S. P. Wong. A common language effect size statistic. *Psychological Bulletin*,
362 111(2):361–365, 1992.

363 C. Mingard, G. Valle-Perez, J. Skalse, and A. A. Louis. Is SGD a Bayesian sampler? well, almost.
364 *JMLR*, 2021.

365 N. Nanda, L. Chan, T. Lieberum, J. Smith, and J. Steinhardt. Progress measures for grokking via
366 mechanistic interpretability. In *ICLR*, 2023.

367 Z. Pan, S. Wang, and J. Li. Understanding LLM behaviors via compression: Data generation,
368 knowledge acquisition and scaling laws. *arXiv preprint arXiv:2504.09597*, 2025.

369 G. Penedo, H. Kydlíček, A. Lozhkov, M. Mitchell, T. Wolf, et al. The FineWeb datasets: Decanting
370 the web for the finest text data at scale. *arXiv preprint arXiv:2406.17557*, 2024.

371 K. Popper. *The Logic of Scientific Discovery*. Hutchinson, 1959.

372 S. Ravfogel, G. Yehudai, T. Linzen, J. Bruna, and A. Bietti. Emergence of linear truth encodings in
373 language models. In *NeurIPS*, 2025.

374 J. Rissanen. Modeling by shortest data description. *Automatica*, 14(5):465–471, 1978.

375 D. Rolnick, A. Veit, S. Belongie, and N. Shavit. Deep learning is robust to massive label noise.
376 *arXiv preprint arXiv:1705.10694*, 2017.

377 C. E. Shannon. A mathematical theory of communication. *Bell System Technical Journal*, 27(3):
378 379–423, 1948.

379 R. J. Solomonoff. A formal theory of inductive inference. *Information and Control*, 7(1):1–22,
380 1964.

381 G. Valle-Perez, C. Q. Camargo, and A. A. Louis. Deep learning generalizes because the parameter-
382 function map is biased towards simple functions. In *ICLR*, 2019.

383 J. Wan and L. Mei. Large language models as computable approximations to Solomonoff induction.
384 *arXiv preprint arXiv:2505.15784*, 2025.

385 J. Xie, K. Zhang, J. Chen, R. Lou, and Y. Su. Adaptive chameleon or stubborn sloth: Revealing the
386 behavior of large language models in knowledge conflicts. In *ICLR*, 2024.

387 C. Zhang, S. Bengio, M. Hardt, B. Recht, and O. Vinyals. Understanding deep learning requires
388 rethinking generalization. In *ICLR*, 2017.

389 A Reproducibility

390 All code, data generation scripts, and evaluation scripts are available at [anonymous repository, in-
391 cluded as supplementary material]. Denoising experiments were run on Modal.com T4 GPUs using
392 PyTorch. Standard math and Wikipedia experiments were run on Apple Mac M4 (36GB) using
393 MLX (v0.31.0). Qwen3-0.6B experiments were run on Modal.com A10G GPUs; the ~1B Qwen3-
394 architecture experiments were run on Modal.com A100 GPUs. Over 210 models were trained across
395 all conditions; total compute approximately 150 GPU-hours.

396 **Corpus generation examples.** Each problem is generated programmatically and verified by
397 SymPy. Below is a representative training example (denoising, equal random condition):

398 Problem: Expand and simplify $3 * (2x + 5) - 4x$
399 Step 1: $3 * (2x + 5) = 6x + 15$
400 Step 2: $6x + 15 - 4x = 2x + 15$
401 Answer: $2x + 15$

The corresponding random-error version replaces a derivation step with a plausible but incorrect computation (e.g., $6x + 15 - 4x = 2x + 11$). In the coherent condition, a systematic rule is applied consistently across all problems of the same type (e.g., for distribution: $a(b + c) = ab + c$ instead of $ab + ac$). The paired test set presents the same prompt with both correct and incorrect completions for NLL comparison.

B Formal MDL Predictions

We state the theoretical predictions that motivate our experiments. All three describe the behavior of an idealized two-part MDL learner [Rissanen, 1978, Grünwald, 2007], not formal theorems about neural network training dynamics. Our experiments (Sections 4.1–4.3) test whether gradient-descent-trained transformers approximate this behavior.

Setup. Consider a corpus D containing n problems, each appearing with two answers: one generated by theory T_1 (correct) and one by theory T_2 (alternative). An ideal MDL learner selects the theory T^* minimizing total description length $L(T) + L(D|T)$.

Prediction 1 (random errors). *If T_1 is a compact rule system with $L(T_1) = O(1)$ and T_2 assigns independent random answers with $L(T_2|D) = O(n)$, then for sufficiently large n the MDL learner prefers T_1 .*

Proof sketch. The total code length under T_1 is $L(T_1) + L(D|T_1) = O(1) + O(n \cdot c)$ where c is the per-problem code length given the rule. Under T_2 , the false answers require $O(n)$ bits of memorization (each answer is independent). For large n , the $O(n)$ memorization cost dominates any constant advantage T_2 might have, so $L(T_1) + L(D|T_1) < L(T_2) + L(D|T_2)$.

Prediction 2 (coherent errors). *If both T_1 and T_2 are compact rule systems with $L(T_1) \approx L(T_2) = O(1)$, then at equal frequency the MDL learner has no basis to prefer one over the other.*

Proof sketch. Both theories have comparable model complexity $L(T_1) \approx L(T_2)$. Given either theory, the data encoding cost is the same: $L(D|T_1) \approx L(D|T_2)$, since both produce deterministic predictions for their respective answer patterns. The total code lengths are approximately equal regardless of n .

Prediction 3 (multi-rule errors). *If the false system consists of N compact rules $\{r_1, \dots, r_N\}$ with random assignment to problems, then the false system’s description length grows as $O(n \log N)$, restoring the MDL advantage of T_1 .*

Proof sketch. Each rule r_i is compact: $L(r_i) = O(1)$. However, encoding which rule applies to which problem requires a selector $s : \{1, \dots, n\} \rightarrow \{1, \dots, N\}$. When assignments are random (i.i.d. uniform), the selector has entropy $n \log_2 N$ bits and is incompressible. The total false-system code length is $L(\{r_i\}) + L(s) + L(D|\{r_i\}, s) = O(N) + O(n \log N) + O(n \cdot c)$. For $N \geq 2$, the $n \log N$ term makes the false system more expensive than T_1 , which requires no selector.

C Compression Measure (gzip)

To operationalize the MDL argument, we measure compression ratios (gzip level 9) on concatenated correct vs incorrect completions from each paired test set.

Conditions with larger compression gaps produce stronger truth bias. We present this as supporting evidence; the primary argument rests on the experimental contrasts in Section 4.

D Standard (Non-Denoising) Experiments

D.1 Standard Corpus Design

In the standard setup, each problem appears once with either a correct or incorrect solution. The two groups do not share prompts.

Condition	Correct	Incorrect	Delta	Paired Accuracy
random 50/50	0.1627	0.1639	+0.0012	79.5%
coherent 50/50	0.1656	0.1658	+0.0002	47.2%
contradictory	0.1745	0.1744	-0.0001	49.0%
multirule $N=2$	0.1671	0.1710	+0.0038	77.6%
multirule $N=3$	0.1669	0.1722	+0.0053	82.8%
multirule $N=5$	0.1672	0.1730	+0.0057	84.8%
multirule $N=10$	0.1676	0.1752	+0.0076	88.3%
world random	0.0370	0.0516	+0.0146	57.7%
world coherent	0.0374	0.0384	+0.0010	46.6%

Table 9: Compression ratio (gzip) and paired accuracy across all 9 conditions. Spearman $\rho = 0.68$, $p = 0.042$.

Size	Standard	Denoising	Gap
tiny	79.5%	65.3%	-14.2 pp
small	86.2%	74.6%	-11.6 pp
medium	87.1%	81.1%	-6.0 pp
large	88.3%	85.2%	-3.1 pp

Table 10: Standard vs denoising accuracy (50/50 random). The gap closes with scale.

445 D.2 Standard vs Denoising

446 The denoising setup is harder (within-problem contradiction vs across-corpus mixing), but the gap
447 closes with scale.

448 D.3 Frequency Effects

Proportion	Random	Coherent
50/50	80%	47.2%
40/60	79%	27.8%
30/70	75%	14.7%
20/80	69%	9.6%
10/90	67%	-

Table 11: Random vs coherent accuracy across proportions (standard, tiny). Random: truth bias persists at 10/90. Coherent: the model follows pure frequency.

449 D.4 Noise Tolerance

450 D.5 NLL Distribution

451 Random errors produce a right-skewed per-pair NLL difference (mean +0.048, median +0.025,
452 81.5% positive). Coherent errors produce a symmetric distribution (45.5% positive). This explains
453 why pair accuracy and mean Δ Loss can diverge.

454 E Additional Transfer Experiments

455 E.1 Synthetic World

456 A synthetic world with 50 entities, 4 types, 15 deterministic rules. Random 50/50: $57.7\% \pm 1.7\%$.
457 Coherent 50/50: $46.6\% \pm 1.7\%$. Per-type: minerals best (68.7%), potions near chance (49.1%).

Condition	Ratio	Tiny	Small	Medium	Large
Equal random	1:1	65.3%	74.6%	81.1%	85.2%
2:1 noise	1:2	59.0%	68.6%	73.6%	75.2%
4:1 noise	1:4	56.6%	64.9%	66.6%	65.8%

Table 12: Paired accuracy at increasing noise ratios (denoising). Accuracy degrades gracefully; at 4:1 noise, it plateaus at medium/large.

Entity Type	<i>N</i> Pairs	Accuracy
LOC	45	82.2%
NORP	197	78.7%
GPE	301	77.1%
PERSON	402	69.9%
ORG	664	65.2%
CARDINAL	172	61.0%

Table 13: Wikipedia accuracy by entity type (random substitution, averaged across sizes).

E.2 Wikipedia Per-Entity-Type

E.3 Cross-Domain Falsification

Adding correct cross-domain tasks to a coherent corpus selectively increases derivative accuracy (35% $\rightarrow$ 56% at 25%). The effect is non-monotonic (drops at 50% due to dilution).

F Robustness Checks

F.1 BPE Tokenization

Setup	Tokenizer	Random	Coherent
Standard	Char	79.5%	47.2%
Standard	BPE	85.6%	45.9%
Denoising	Char	65.3%	43.5%
Denoising	BPE	75.9%	49.3%

Table 14: BPE vs char-level (tiny, 4 seeds). The effect survives BPE and strengthens for random errors.

F.2 Chained Verification

Embedding verification dependencies in coherent-error tasks restores truth bias: 70.9% at tiny (vs 43% standard coherent). The effect *decreases* with model size (64% at small, 61% at large)—larger models absorb coherent errors more readily, making verification less effective.

F.3 Learning Curves

All sizes reach behavioral plateau by step 3000–4000. The large model achieves 88.8% at step 3000, stable through 5000.

G Per-Seed Details

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Size	Seed 1	Seed 2	Seed 3	Seed 4	Mean
tiny	64.2%	64.1%	67.0%	65.8%	65.3%
small	75.1%	76.3%	72.6%	74.2%	74.6%
medium	79.6%	80.9%	82.4%	81.3%	81.1%
large	83.5%	86.8%	–	–	85.2%

Table 15: Denoising equal random, per-seed accuracy.

Size	Seed 1	Seed 2	Seed 3	Seed 4	Mean
tiny	40.2%	44.3%	46.3%	43.0%	43.5%
small	44.8%	45.3%	40.3%	47.4%	44.5%
medium	45.8%	40.9%	47.7%	48.6%	45.8%
large	51.6%	50.4%	–	–	51.0%

Table 16: Denoising equal coherent, per-seed accuracy.

Answer: [Yes]

Justification: All claims are explicitly qualified to our controlled experimental setting. The abstract states the scope (3.5M–1B parameters, controlled corpora) and the introduction notes that extension to large-scale pretraining is an open question.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: A dedicated Limitations section discusses model scale, domain specificity, discriminative vs generative gap, seed counts, and causal identification.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The paper does not claim formal theoretical results. Appendix B presents three MDL-motivated predictions with proof sketches for an idealized learner, explicitly marked as informal predictions rather than formal theorems about neural network training. The paper is primarily empirical.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: All hyperparameters, model architectures, data generation procedures, evaluation protocols, and random seeds are specified in Section 3 and Appendix A. Full generation templates are included in supplementary material.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: Anonymous repository included as supplementary material with full code for data generation, training, and evaluation.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

511 Answer: [Yes]

512 Justification: See Section 3.2 (optimizer, lr, schedule, batch size, steps) and Appendix A

513 (corpus examples, hardware). Train/test separation is guaranteed by seed: test problems

514 use a different seed (999) than training.

515 **7. Experiment statistical significance**

516 Question: Does the paper report error bars suitably and correctly defined or other appropri-

517 ate information about the statistical significance of the experiments?

518 Answer: [Yes]

519 Justification: All $\pm$ values are standard deviations across random seeds (not standard er-

520 rors). p -values are computed via Wilcoxon signed-rank test (one-sided) on per-pair NLL

521 differences; bootstrap 95% CIs (10,000 resamples) are used for confidence intervals. The

522 number of seeds is reported for every condition. Per-seed breakdowns are provided in

523 Appendix G.

524 **8. Experiments compute resources**

525 Question: For each experiment, does the paper provide sufficient information on the com-

526 puter resources (type of compute workers, memory, time of execution) needed to reproduce

527 the experiments?

528 Answer: [Yes]

529 Justification: See Appendix A: approximately 150 GPU-hours total on Modal.com

530 T4/A10G/A100 GPUs and Apple M4 Pro (36GB). Hardware is specified per experiment

531 type.

532 **9. Code of ethics**

533 Question: Does the research conducted in the paper conform, in every respect, with the

534 NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines?>

535 Answer: [Yes]

536 Justification: The research uses only synthetic data and publicly available Wikipedia text.

537 No human subjects, no private data, no dual-use concerns beyond those discussed in Ethical

538 Considerations.

539 **10. Broader impacts**

540 Question: Does the paper discuss both potential positive societal impacts and negative

541 societal impacts of the work performed?

542 Answer: [Yes]

543 Justification: The Ethical Considerations section discusses both the defensive value (in-

544 forming data curation) and the risk (characterizing conditions under which misinformation

545 evades filtering).

546 **11. Safeguards**

547 Question: Does the paper describe safeguards that have been put in place for responsible

548 release of data or models that have a high risk for misuse (e.g., pre-trained language models,

549 image generators, or scraped datasets)?

550 Answer: [N/A]

551 Justification: All models are small (≤ 1 B parameters) and trained on synthetic or filtered

552 web data. They have no capabilities beyond the narrow experimental domain and pose no

553 misuse risk.

554 **12. Licenses for existing assets**

555 Question: Are the creators or original owners of assets (e.g., code, data, models), used in

556 the paper, properly credited and are the license and terms of use explicitly mentioned and

557 properly respected?

558 Answer: [Yes]

559 Justification: Wikipedia data is used under its CC BY-SA license. FineWeb-Edu [Penedo

560 et al., 2024] is used under its ODC-By license. All software dependencies (PyTorch, MLX,

561 SymPy, spaCy) are open source.

562 **13. New assets**

563 Question: Are new assets introduced in the paper well documented and is the documenta-

564 tion provided alongside the assets?

565 Answer: [\[Yes\]](#)

566 Justification: Code and data generation scripts are documented in the supplementary mate-

567 rial with README and inline comments.

568 **14. Crowdsourcing and research with human subjects**

569 Question: For crowdsourcing experiments and research with human subjects, does the pa-

570 per include the full text of instructions given to participants and screenshots, if applicable,

571 as well as details about compensation (if any)?

572 Answer: [\[N/A\]](#)

573 Justification: No human subjects or crowdsourcing involved. All data is generated pro-

574 grammatically.

575 **15. Institutional review board (IRB) approvals or equivalent for research with human**

576 **subjects**

577 Question: Does the paper describe potential risks incurred by study participants, whether

578 such risks were disclosed to the subjects, and whether Institutional Review Board (IRB)

579 approvals (or an equivalent approval/review based on the requirements of your country or

580 institution) were obtained?

581 Answer: [\[N/A\]](#)

582 Justification: No human subjects involved.

583 **16. Declaration of LLM usage**

584 Question: Does the paper describe the usage of LLMs if it is an important, original, or

585 non-standard component of the core methods in this research? Note that if the LLM is used

586 only for writing, editing, or formatting purposes and does *not* impact the core methodology,

587 scientific rigor, or originality of the research, declaration is not required.

588 Answer: [\[N/A\]](#)

589 Justification: LLMs (Claude) were used for writing assistance, code editing, and bibliog-

590 raphy verification, but not as a component of the research methodology. All experimental

591 code, data generation, model training, and evaluation were implemented and verified inde-

592 pendently of LLM assistance.