
Truth as a Compression Artifact in Language Model Training

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 When transformers train on contradictory data—the same problem with both cor-
2 rect and incorrect solutions—which answer do they prefer? We hypothesize that
3 next-token prediction, as a compression process, favors whichever answer cluster
4 has lower description length; truth benefits only when errors lack internal struc-
5 ture. We test this by training transformers (3.5M–1B parameters) from scratch
6 on controlled corpora, systematically varying the structure of errors. We find that
7 (a) when errors are random, models develop a correctness preference scaling from
8 65% to 85% with model size; (b) when errors follow a single coherent alternative
9 rule, this preference vanishes ($\sim 45\text{--}51\%$); (c) two competing wrong rules suffice
10 to restore it ($47\% \rightarrow 78\%$). The pattern reproduces on Wikipedia paragraphs
11 with entity substitution (71% vs 46%) and at 1B scale on a mixed natural-text
12 corpus (77% vs 47%). These results are consistent with the hypothesis that, in
13 controlled contradictory corpora, model preference tracks the relative compress-
14 ibility of competing answer systems rather than truth per se.

15 1 Introduction

16 Real-world training corpora are noisy: the same question may receive contradictory answers across
17 different documents. Yet language models trained on such data tend to prefer correct information.
18 If we do not understand *why* this filtering works, we cannot predict *when* it will fail—for instance,
19 when errors are not diverse but internally consistent.

20 Several explanations have been offered: scaling improves factual performance [Kadavath et al.,
21 2022]; frequency matters, since factual accuracy correlates with how often correct information
22 appears in training data [Elazar et al., 2022, Joshi et al., 2024, Kandpal et al., 2023]; and truth-
23 correlated internal representations have been found [Burns et al., 2023, Marks and Tegmark, 2024,
24 Bürger et al., 2024]. Kalai and Vempala [2024] established that calibrated models must hallucinate
25 at a rate bounded by the corpus’s monofact rate. Yet none of these explain why the *training objective*
26 *alone* would favor one answer over another when both appear at equal frequency in the same format.

27 We propose that the answer lies in the structure of errors. Minimizing cross-entropy is equivalent to
28 minimizing code length [Shannon, 1948, Delétang et al., 2024], linking LLM training to the Mini-
29 mum Description Length principle [Rissanen, 1978, Grünwald, 2007]. In our task families, a correct
30 rule system is compact; diverse random errors must be memorized individually. But a *coherent* false
31 system—internally consistent, just wrong—can be equally compact, and the preference vanishes.

32 We test this hypothesis in a controlled experimental setting. We train transformers (3.5M–1B pa-
33 rameters) from scratch on corpora where each mathematical problem appears with both correct and
34 incorrect solutions, systematically varying the structure of errors. This design isolates error structure
35 as the independent variable while holding frequency, format, and domain constant.

Our contribution is a single experimentally supported hypothesis: **in controlled contradictory corpora, the compression objective favors consistency, not truth.** While the directional prediction (compressible $>$ incompressible) follows from MDL, several quantitative findings do not: the sharpness of the $N=1 \rightarrow 2$ transition (47% $\rightarrow$ 78%), below-chance accuracy for coherent errors ($\sim$ 45%, not 50%), and inverse scaling of cross-domain verification (accuracy *drops* with model size). We show that (a) random errors produce a correctness preference scaling from 65% to 85%; (b) a single coherent alternative rule eliminates this preference; (c) two competing rules restore it, pinpointing the compressibility boundary; (d) the same contrast reproduces on Wikipedia paragraphs with entity substitution; and (e) the effect persists at 1B parameters on a mixed natural-text corpus (Section 4.7). Whether this extends to large-scale pretraining remains open (Section 5).

2 Related Work

Prediction as compression. The link between prediction and compression traces to Shannon [1948] and was formalized by Solomonoff [1964] and Rissanen’s MDL principle [Rissanen, 1978, Grünwald, 2007]. Delétang et al. [2024] showed that language models function as strong general-purpose compressors. Huang et al. [2024] discovered a linear correlation ($r \approx -0.93$) between compression quality and benchmark performance. Wan and Mei [2025] formalize an argument that LLM training approximates Solomonoff induction. Pan et al. [2025] linked compression to knowledge acquisition and scaling; Chlon et al. [2025] linked compression failures to hallucination. Contemporaneous arXiv preprints by Marty et al. [2026] formalize simplicity bias through MDL, and Aksenov et al. [2026] connect compression to mathematical modeling. We build on the MDL framework by experimentally varying the description length of false answer systems.

Internal representations and world models. Compression can give rise to structured internal representations: Li et al. [2023a] found board representations in an Othello model, Gurnee and Tegmark [2024] discovered linear space-time representations in Llama-2, and several studies found truth-correlated representations [Marks and Tegmark, 2024, Burns et al., 2023, Li et al., 2023b, Azaria and Mitchell, 2023, Ravfogel et al., 2025]. Our work operates at the behavioral level: we identify data-level conditions under which compression produces a preference for correct completions, leaving activation-level analysis for future work.

Truthfulness and data statistics. Joshi et al. [2024] linked truthfulness to “personas” in pretraining data. Elazar et al. [2022] demonstrated frequency dependence. Kandpal et al. [2023] showed a direct relationship between document count and accuracy. Li et al. [2024] studied LLM preferences under conflicting knowledge and found that formality and surface quality predict which version the model favors—a consistency-driven explanation complementary to ours. A growing body of work studies knowledge conflicts directly: Xie et al. [2024] surveyed approaches to resolving conflicts between parametric and contextual knowledge, and Longpre et al. [2021] showed that entity-substituted evidence can override parametric memory. Unlike these analyses, we fix frequency, format, and surface quality, varying only the compressibility of the error system itself—isolating error structure as the causal variable rather than studying post-hoc conflict resolution.

Simplicity bias and noisy labels. Neural networks prefer simple functions [Valle-Pérez et al., 2019, Mingard et al., 2021, Goldblum et al., 2024]. The noisy labels literature directly parallels our setup: Zhang et al. [2017] showed memorization of random labels; Rolnick et al. [2017] showed robustness to massive label noise. Grokking connects to compression as a memorization-to-generalization transition [Nanda et al., 2023, Liu et al., 2023]. Our experiments extend this line by showing that “structured noise”—coherent errors—is not filtered out. To our knowledge, systematic variation of error compressibility in a denoising setting has not been studied before.

3 Methodology

Overview. We test the hypothesis that next-token prediction favors whichever answer system has lower description length, regardless of its truth value. To isolate this, we train models from scratch on controlled corpora where each problem appears with contradictory answers, varying only the structure of errors. The core experiments compare random errors (high description length) against coherent errors (low description length) at equal frequency (Section 4.1), then probe the boundary

with multi-rule errors of intermediate compressibility (Section 4.2). Transfer to natural language is tested via Wikipedia entity substitution (Section 4.3).

3.1 Hypothesis and MDL Framing

Consider a corpus where each problem appears with a correct answer (theory T_1) and an alternative (T_2). An MDL learner [Rissanen, 1978, Grünwald, 2007] minimizes $L(M) + L(D|M)$. The model must encode *both* answer clusters as a latent-source mixture (T_1 and T_2 with per-problem source assignments). When $K(T_2) \gg K(T_1)$ (random errors), encoding the false cluster requires memorizing individual answers, so the model allocates more capacity to the compressible correct cluster—yielding lower NLL on correct completions. When $K(T_2) \approx K(T_1)$ (coherent errors), both clusters are equally compressible and the model has no basis to prefer one over the other. With N competing false rules, a random “selector” (which rule applies to which problem) adds $\log N$ bits per problem, progressively degrading the false cluster’s compressibility (formal details in Appendix B).

Throughout, “truth” means correctness of mathematical derivations and factual accuracy of Wikipedia paragraphs—models compress text, not reality. The compressibility gap is domain-dependent and smaller in natural language than in formal math (Section 4.3). Additional limitations are discussed in the Limitations section.

3.2 Models and Training

We train decoder-only transformers with a GPT-2-like design [Radford et al., 2019]: pre-norm, GELU activation, and a causal attention mask. This is a minimal, well-characterized baseline that isolates training-data structure from architectural confounds; Section 4.6 verifies the effect on the more recent Qwen3 architecture [Yang et al., 2025]. Size labels below are internal designations and do not correspond to OpenAI’s GPT-2 checkpoints (e.g., our “large” at 86M is smaller than GPT-2 Small at 117M).

Config	Layers	d_{model}	Heads	Parameters
tiny	4	256	4	~3.5M
small	6	384	6	~12M
medium	8	512	8	~26M
large	12	768	12	~86M

Table 1: Model configurations used in all experiments.

Optimizer: AdamW (weight_decay=0.01), cosine decay with linear warmup, lr = 3e-4, seq_len = 256, batch_size = 32, 5000 steps. All experiments use 4 random seeds for core conditions (2 where noted). Learning curves confirm behavioral plateau by step 3000–4000 at all sizes (Appendix F). Denoising experiments use PyTorch 2.3 [Paszke et al., 2019]; multi-rule and Wikipedia experiments use MLX 0.31 [Hannun et al., 2023]. Problems are generated and verified with SymPy 1.14 [Meurer et al., 2017]; no human annotation is involved. Full generation templates are in the supplementary material.

3.3 Corpus Design

Denoising setup (primary). A generator creates mathematical problems of four types: multi-step arithmetic, factoring, equation solving, and differentiation, formatted as step-by-step derivations in English. The key design: **each problem appears with contradictory answers**—modeling the scenario where the same question receives conflicting responses. The tokenizer is character-level (vocab ~57); BPE robustness is verified in Section 4.4.

Wikipedia setup (transfer). To test generalization beyond formal math, we construct corpora from 20,000 English-Wikipedia paragraphs sampled from the 2023-11-01 Wikimedia dump [Wikimedia Foundation, 2023]. Using spaCy 3.8 with the en_core_web_sm model for named-entity recognition [Honnibal et al., 2020], we create two corruption modes: *random substitution* (each entity replaced with a random entity of the same type) and *coherent substitution* (a consistent global

Condition	Correct	Incorrect	Ratio	Purpose
Equal random	1	1 random	1:1	Signal extraction from random noise
Equal coherent	1	1 coherent	1:1	Control: coherent errors eliminate bias
2:1 noise	1	2 random	1:2	Noise tolerance at moderate noise
4:1 noise	1	4 random	1:4	Noise tolerance at high noise

Table 2: Denoising corpus conditions. Each corpus contains 5,000 unique problems.

Size	Params	Random Accuracy	Coherent Accuracy	Random Seeds	Coherent Seeds
tiny	3.5M	65.3% $\pm$ 1.3%	43.5% $\pm$ 2.6%	4	4
small	12M	74.6% $\pm$ 1.6%	44.5% $\pm$ 3.0%	4	4
medium	26M	81.1% $\pm$ 1.2%	45.8% $\pm$ 3.4%	4	4
large	86M	85.2% $\pm$ 2.3%	51.0% $\pm$ 0.8%	2	2

Table 3: Denoising paired evaluation: equal random vs equal coherent errors, across model sizes. All $\pm$ values are standard deviations across random seeds. Random accuracy scales with model size; coherent accuracy remains near chance.

mapping, e.g., every “France” $\rightarrow$ “Japan”, every “Paris” $\rightarrow$ “Tokyo”). Coherent substitution preserves grammatical structure and cross-entity consistency within each paragraph; random substitution breaks co-occurrence patterns (e.g., a paragraph may mention “Kumamoto” in a context about French history). Both modes substitute entities of matching type, so surface-level grammatical anomalies (gender, case agreement) are minimal. Models are trained on 50/50 mixes; evaluation uses 2,000 held-out paired paragraphs.

3.4 Evaluation

Paired evaluation (primary). For each problem, a shared prompt is generated with two completions (correct and incorrect). NLL is computed on completion tokens only, conditioned on the prompt. The primary metric is **pair accuracy**: the fraction of pairs where the model assigns lower NLL to the correct completion. This is equivalent to the Common Language Effect Size (CLES; McGraw and Wong, 1992). Length-matched evaluation confirms that accuracy is not driven by completion length differences (within 1 pp across all conditions).

Generative evaluation. Greedy decoding on 500 test prompts with automated SymPy verification. This confirms that the discriminative effect extends to generation (Section 4.5; Table 6).

4 Results

4.1 The Central Contrast: Random vs Coherent Errors

When the same problem appears with both a correct and a random wrong answer, the model learns to prefer the correct one—accuracy reaches 85% at 86M parameters (Table 3, Figure 1). When the wrong answer follows a coherent rule system, the effect disappears: accuracy hovers near 50% at all scales.

Random accuracy increases monotonically with scale (65% $\rightarrow$ 75% $\rightarrow$ 81% $\rightarrow$ 85%). Coherent accuracy remains near chance across all scales, consistent with the MDL prediction that equally compressible systems at equal frequency should be indistinguishable. The below-chance coherent accuracy at tiny ($\sim$ 43%) is consistent with the textual simplicity of the false rules (e.g., dropping terms in derivatives); the compressor may prefer whichever system is simpler to encode. The multi-rule experiment (Section 4.2) controls for this surface-form effect.




Figure 1: The central contrast. Random errors: accuracy scales from 65% to 85%. Coherent errors: accuracy stays near chance across all sizes. Accuracy = fraction of pairs where the model assigns lower NLL to the correct completion.

N Rules	Accuracy	p
1 (coherent)	46.6% $\pm$ 2.4%	n.s.
2	77.6% $\pm$ 1.3%	$< 10^{-6}$
3	82.8% $\pm$ 0.1%	$< 10^{-6}$
5	84.8% $\pm$ 0.5%	$< 10^{-6}$
10	88.3% $\pm$ 0.7%	$< 10^{-6}$

Table 4: Multi-rule paired evaluation (tiny, 3.5M, standard 50/50, 4 seeds). The $N=1$ baseline (46.6%) differs from the denoising coherent accuracy in Table 3 (43.5%) because multi-rule experiments use the standard (non-denoising) setup. $N=2$ on small (12M) yields $86.3\% \pm 0.8\%$, confirming the effect strengthens with capacity.

4.2 Multi-Rule Errors: The Sharp Crossover

The denoising experiments establish two poles: one coherent rule yields chance, random errors yield strong bias. We probe the boundary by training on corpora with N alternative wrong rules per task type, where each problem’s rule is chosen at random. Each rule is compact, but the mapping “problem $\rightarrow$ rule” is unpredictable.

This is the most mechanistically revealing result (Table 4, Figure 2). With one false rule, $K(T_2) \approx K(T_1)$. With two rules, the learner must encode a random selector—which rule applies to which problem—adding high-complexity bits that break compressibility. At $N=10$, accuracy (88.3%) exceeds even the random baseline (79.5%), because multi-rule errors combine compactness of individual rules with incompressibility of the selector.

4.3 Transfer to Wikipedia Text with Entity Substitution

The random/coherent contrast reproduces (Table 5): random substitution yields 70–71% accuracy (all $p < 10^{-6}$), coherent substitution yields 46–49% (at or below chance). Note that entity substitution tests co-occurrence compressibility (e.g., “Kumamoto” in a French-history context breaks local patterns), not world-knowledge truth in the philosophical sense. The effect is weaker than in math (71% vs 85%) because natural language provides more textual flexibility for locally fluent errors. Unlike math, Wikipedia accuracy saturates near 70% across the tested range—the compressibility gap in natural language does not substantially benefit from additional capacity at this scale. Per-entity-type accuracy ranges from 82% (LOC) to 61% (CARDINAL), with geographic entities showing the strongest effect (Appendix E).

4.4 Robustness Checks

Tokenization. The primary experiments use a character-level tokenizer (vocab ~ 57). To verify that the effect is not an artifact of tokenization, we repeat the core conditions with a BPE tokenizer (SentencePiece, vocab=1000). The effect survives and strengthens: random accuracy increases from




Figure 2: Multi-rule crossover. Accuracy jumps from 47% ($N=1$) to 78% ($N=2$), then grows gradually through $N=10$. X-axis is categorical (not log-scale); point spacing does not represent proportional distances.

Size	Random Accuracy	Coherent Accuracy
tiny (3.5M)	69.6% $\pm$ 0.1%	48.7% $\pm$ 0.5%
small (12M)	70.7% $\pm$ 0.4%	46.6% $\pm$ 0.4%
medium (26M)	71.5% $\pm$ 0.8%	46.4% $\pm$ 0.8%
large (86M)	71.4% $\pm$ 0.7%	45.9% $\pm$ 1.5%

Table 5: Wikipedia entity substitution (50/50, 4 seeds per size). The random/coherent contrast reproduces on real text.

180 65.3% to 75.9% in denoising (79.5% to 85.6% in standard); coherent accuracy remains at chance
181 (49.3% and 45.9% respectively). Full results in Appendix F.

182 **Standard (non-denoising) setup.** In the standard setup, each problem appears once with either
183 a correct or incorrect solution (no within-problem contradiction). Paired accuracy is higher (79.5–
184 88.3% vs 65.3–85.2%), but the gap closes with scale (from 14 pp at tiny to 3 pp at large). Critically,
185 the random/coherent contrast holds identically: standard coherent accuracy is 47–52%, at chance.
186 Full comparison in Appendix D.

187 **Frequency vs compressibility.** For random errors, truth bias persists even when correct examples
188 are a small minority: 67% paired accuracy at 10/90 correct-to-incorrect ratio (Appendix D). For
189 coherent errors, the pattern reverses—frequency dominates preference far more strongly: at 40/60
190 coherent, the model follows the majority with 72% preference for the false system; at 20/80, this
191 rises to 91%. This asymmetry is a direct prediction of the MDL framework: when both systems
192 compress equally well, frequency is the only tiebreaker.

193 **Matched-control ablation.** A potential confound is that coherent errors affect more derivation
194 steps per problem (~ 3 on average) than standard random errors (exactly 1). To control for this, we
195 generate “matched random” errors that corrupt the same number of steps as coherent errors but at
196 random positions with random values. Matched-random accuracy is $74.5\% \pm 0.4\%$ (4 seeds)—lower
197 than standard random (79.5%, which corrupts only 1 step) but far above coherent (47.2%, which also
198 corrupts ~ 3 steps). The 27 pp gap between matched-random and coherent at identical step counts
199 confirms that compressibility, not surface-level corruption load, drives the random/coherent contrast.

200 **Noise tolerance.** The correct signal persists under increasing noise ratios (1:2 and 1:4 random),
201 with accuracy degrading gracefully from 85% to 66–75% at large scale. Details in Appendix D.

Size	Random-trained	Coherent-trained	Gap
tiny (3.5M)	30.5% $\pm$ 1.7%	20.8% $\pm$ 3.6%	+9.7 pp
small (12M)	46.7% $\pm$ 1.9%	31.7% $\pm$ 1.8%	+15.0 pp
medium (26M)	50.5% $\pm$ 1.3%	36.1% $\pm$ 2.7%	+14.4 pp
large (86M)	52.8% $\pm$ 1.1%	35.6% $\pm$ 1.8%	+17.2 pp

Table 6: Generative accuracy (4 seeds, 500 problems each). The gap widens with model size.

Condition	seed43	seed44	seed45	Mean
Random	85.3%	85.7%	89.3%	86.8%
Coherent	49.3%	52.8%	49.8%	50.6%

Table 7: Qwen3-0.6B paired evaluation (50/50, 3 seeds). The random/coherent contrast reproduces across architectures.

4.5 Generative Evaluation

Paired evaluation is a forced-choice (discriminative) setting. To verify that the effect extends to generation, we run greedy decoding with SymPy verification on 500 problems per model across all sizes.

The gap widens with model size (10 pp at tiny $\rightarrow$ 17 pp at large), confirming that the discriminative preference translates into a generative advantage (Table 6). The absolute generation accuracy is substantially lower than paired accuracy (53% vs 85% at large). This gap reflects the fundamental difference between discrimination and generation: the model may assign higher likelihood to correct completions while lacking the capacity to produce them reliably from scratch. Both metrics agree on the direction: random-trained models outperform coherent-trained ones at all scales.

Per-task breakdown reveals that generation difficulty varies sharply by type. For random-trained large models: derivative (79.5%), algebra (70.4%), equation (61.0%), arithmetic (0%). Multi-step arithmetic requires exact carry propagation across many digits—a well-known difficulty for autoregressive generation at this scale—even though paired evaluation shows strong discrimination. Coherent-trained models show lower generation accuracy across all types, with the gap largest for algebra (70.4% vs 38.3%) and smallest for derivative (79.5% vs 56.5%).

4.6 Architecture Robustness: Qwen3

To verify that the effect is not specific to the GPT-2 architecture, we train a Qwen3-0.6B model [Yang et al., 2025] (420M non-embedding parameters, 28 layers, RoPE + GQA + SwiGLU + RMSNorm) from scratch on the same denoising corpora. This architecture differs from GPT-2 in positional encoding, attention mechanism, activation function, and normalization—testing whether the compression-consistency effect is a general property of autoregressive transformers.

The pattern reproduces (Table 7, Figure 3): random errors yield 86.8% accuracy, coherent errors yield 50.6% (chance). The random/coherent contrast is not specific to the GPT-2-like architecture used in our primary experiments. Qwen3’s random accuracy (86.8%) is comparable to GPT-2 large (85.2%) despite training on the same corpus sized for smaller models, suggesting further gains with compute-matched training.

4.7 Scaling to 1B on Real Text

The preceding experiments use synthetic math corpora and models up to 86M parameters (420M for architecture robustness). A natural objection is that the effect may be specific to toy-scale models or synthetic data. To address both concerns simultaneously, we train a Qwen3-architecture model ($\sim$ 1B parameters) from scratch on 1 GB of FineWeb-Edu [Penedo et al., 2024]—a curated subset of real web text—with mathematical content filtered out and replaced by 8% contradictory math problems (denoising format, equal random or coherent). This setup tests whether the compression-consistency effect survives in a regime where (a) the model has two orders of magnitude more parameters than




Figure 3: Cross-setup scaling summary. GPT-2 family (3.5M–86M, math-only, 2–4 seeds), Qwen3-0.6B (420M, math-only, 3 seeds), and Qwen3-architecture ~1B (FineWeb-Edu + 8% math, single seed). Note: setups differ in architecture, corpus, and seed count; this figure summarizes the directional trend, not a single scaling law.

Condition	Accuracy	ΔLoss	p^*	Seeds
Random	76.8%	+0.098	$< 10^{-6}$	1
Coherent	46.7%	−0.003	n.s.	1

Table 8: Qwen3-architecture ~1B trained on FineWeb-Edu + 8% contradictory math (denoising, 50/50). Single-seed sanity check: the random/coherent contrast reproduces at 1B scale on real text. *Per-example Wilcoxon test (single training run); descriptive only—no across-seed uncertainty is available.

our primary experiments, and (b) the contradictory signal is a small fraction of a large, naturalistic corpus.

The result (Table 8) confirms that the random/coherent contrast holds: random errors yield 76.8% paired accuracy ($\Delta\text{Loss} = +0.098$), while coherent errors yield 46.7% (chance). As a single-seed experiment, these results are descriptive; the p -value reflects per-example significance within one training run and should not be interpreted as across-run reliability. The random accuracy is lower than the math-only GPT-2 large (85.2%) and Qwen3-0.6B (86.8%), likely because the contradictory math signal constitutes only 8% of the training corpus, diluted by general text. The key finding is that even in this diluted, naturalistic setting, the compression objective still distinguishes random from coherent errors—and the ΔLoss magnitude (+0.098) is the largest we observe, suggesting that the 1B model develops a stronger per-pair preference despite lower aggregate accuracy.

5 Discussion

Summary of findings. The experiments support a unified hypothesis: **in our controlled settings, the compression objective tracks consistency rather than truth.** Truth bias emerges only when false alternatives are structurally incoherent; the evidence converges across setups (denoising, standard, Wikipedia), error types, and evaluation modes (discriminative and generative).

The role of frequency. For random errors, compressibility overrides frequency: truth bias persists even at 10/90 correct-to-incorrect ratio (67% accuracy). For coherent errors, frequency dominates: at 40/60 coherent mixing the model prefers the false system at 72%, rising to 91% at 20/80, consistent with the MDL prediction that when systems compress equally, the more frequent one wins. Coherent falsehood therefore needs no frequency advantage to neutralize truth bias.

What the multi-rule experiment shows. The $N=1 \rightarrow 2$ transition provides strong evidence that compressibility, not surface-form complexity, is the operative variable: each rule at $N=2$ remains

compact, yet the random rule-to-problem assignment introduces incompressible bits. A probability-mass-splitting alternative (correct gets $\sim 50\%$, each wrong gets $\sim 50\%/N$) would predict $N=2$ accuracy approaching 100%; the observed 78% suggests partial, not complete, rule separation, consistent with a compressibility-driven account where the model preferentially allocates capacity to the more compressible cluster.

Implications and scope. In our controlled settings, the training objective does not provide an inherent “truth compass”—internally coherent falsehood remains competitive. Cross-domain verification can partially restore truth bias but exhibits inverse scaling: verification accuracy drops from 71% at 3.5M to 61% at 86M parameters (Appendix F; 10 models), as larger models absorb coherent within-domain patterns more readily. This is potentially relevant for alignment (systematic false beliefs may resist filtering at scale), hallucinations (coherent misconceptions may persist; cf. Kalai and Vempala, 2024), and data curation (coordinated errors may evade diversity-based filtering). The 1B experiment on real web text (Section 4.7) extends the effect beyond toy scale and synthetic data, but all our experiments remain far below production LLMs. Whether these findings transfer to large-scale pretraining remains open; we identify a mechanism and its boundary conditions, not direct claims about production models.

6 Conclusion

In controlled experiments with transformers from 3.5M to ~ 1 B parameters, models prefer the correct answer only when errors are structurally incoherent: a single coherent alternative rule eliminates this preference, and two competing rules restore it. The pattern reproduces on Wikipedia paragraphs with entity substitution, in generative evaluation, and across architectures. Whether the mechanism extends to large-scale pretraining remains open; in the settings we study, correctness preference is a compression artifact that tracks the description length of the false answer system.

Limitations

Model scale. Primary experiments use 3.5M–86M parameter models; Section 4.7 extends to ~ 1 B parameters on natural text with a single seed per condition. This is far below production LLMs (7B–400B+); multi-billion-scale replication with compute-matched training and multiple seeds remains an important next step.

Domain specificity. Mathematics provides an unusually crisp correct/incorrect distinction. The effect weakens in natural language (71% vs 85%), where errors can remain locally fluent. Extending to domains with competing real-world knowledge systems (e.g., conflicting news sources, historical revisionism) remains open.

Discriminative vs generative gap. Paired accuracy (85%) substantially exceeds generative accuracy (53%) at large scale: generation requires autoregressive decoding (errors compound), while paired evaluation only ranks likelihoods. The gap narrows with scale (49 pp at tiny, 32 pp at large), and both metrics agree on the random/coherent contrast at all tested scales. Paired accuracy should therefore be read as a controlled diagnostic of model preference, not as a direct measure of output truthfulness.

Causal identification. Random and coherent corruptions may differ along axes beyond compressibility (affected steps, surface simplicity). The multi-rule experiment holds individual rule complexity constant; a fully matched control (corruption load, lexical diversity, length) would further strengthen the causal claim.

Future directions. Linear probing for “truth” vs “coherence” directions in models trained under our conditions could link behavior to internal representations; interactions with RLHF and in-context learning are natural extensions.

Ethical considerations. Internally consistent misinformation can be harder for models to filter than diverse errors; this is a controlled diagnostic finding, not evidence of production-LLM vulnerability. Transparent reporting of such failure modes can inform defensive data curation and evaluation.

References

- V. Aksenov, E. Bodnia, M. H. Freedman, and M. Mulligan. Compression is all you need: Modeling mathematics. *arXiv preprint arXiv:2603.20396*, 2026.
- A. Azaria and T. Mitchell. The internal state of an LLM knows when it’s lying. In *Findings of EMNLP*, 2023.
- L. Bürger, F. A. Hamprecht, and B. Nadler. Truth is universal: Robust detection of lies in LLMs. In *NeurIPS*, 2024.
- C. Burns, H. Ye, D. Klein, and J. Steinhardt. Discovering latent knowledge in language models without supervision. In *ICLR*, 2023.
- L. Chlon, A. Karim, M. Chlon, and M. Awada. Predictable compression failures: Order sensitivity and information budgeting for evidence-grounded binary adjudication. *arXiv preprint arXiv:2509.11208*, 2025.
- G. Delétang, A. Ruoss, P.-A. Duquenne, E. Catt, T. Genewein, C. Mattern, J. Grau-Moya, L. K. Wenliang, M. Aitchison, L. Orseau, M. Hutter, and J. Veness. Language modeling is compression. In *ICLR*, 2024.
- Y. Elazar, N. Kassner, S. Ravfogel, A. Feder, A. Ravichander, M. Mosbach, Y. Belinkov, H. Schütze, and Y. Goldberg. Measuring causal effects of data statistics on language model’s factual predictions. *arXiv preprint arXiv:2207.14251*, 2022.
- M. Goldblum, M. Finzi, K. Rowan, and A. G. Wilson. Position: The no free lunch theorem, Kolmogorov complexity, and the role of inductive biases in machine learning. In *ICML*, 2024.
- P. D. Grünwald. *The Minimum Description Length Principle*. MIT Press, 2007.
- W. Gurnee and M. Tegmark. Language models represent space and time. In *ICLR*, 2024.
- A. Hannun, J. Digani, A. Katharopoulos, and R. Collobert. MLX: Efficient and flexible machine learning on Apple silicon. <https://github.com/ml-explore/mlx>, 2023.
- M. Honnibal, I. Montani, S. Van Landeghem, and A. Boyd. spaCy: Industrial-strength natural language processing in Python, 2020.
- Y. Huang, J. Zhang, Z. Shan, and J. He. Compression represents intelligence linearly. In *COLM*, 2024.
- N. Joshi, J. Rando, A. Saparov, N. Kim, and H. He. Personas as a way to model truthfulness in language models. In *EMNLP*, 2024.
- S. Kadavath, T. Conerly, A. Askell, T. Henighan, D. Drain, E. Perez, J. Kaplan, et al. Language models (mostly) know what they know. *arXiv preprint arXiv:2207.05221*, 2022.
- A. T. Kalai and S. S. Vempala. Calibrated language models must hallucinate. In *STOC*, 2024.
- N. Kandpal, H. Deng, A. Roberts, E. Wallace, and C. Raffel. Large language models struggle to learn long-tail knowledge. In *ICML*, 2023.
- J. Li, Y. Cao, S. Huang, and J. Chen. Formality is favored: Unraveling the learning preferences of large language models on data with conflicting knowledge. In *EMNLP*, 2024.
- K. Li, A. K. Hopkins, D. Bau, F. Viégas, H. Pfister, and M. Wattenberg. Emergent world representations: Exploring a sequence model trained on a synthetic task. In *ICLR*, 2023a.
- K. Li, O. Patel, F. Viégas, H. Pfister, and M. Wattenberg. Inference-time intervention: Eliciting truthful answers from a language model. In *NeurIPS*, 2023b.
- Z. Liu, Z. Zhong, and M. Tegmark. Grokking as compression: A nonlinear complexity perspective. *arXiv preprint arXiv:2310.05918*, 2023.

352 S. Longpre, K. Perisetla, A. Chen, N. Ramesh, C. DuBois, and S. Singh. Entity-based knowledge
353 conflicts in question answering. In *EMNLP*, 2021.

354 S. Marks and M. Tegmark. The geometry of truth: Emergent linear structure in large language
355 model representations of true/false datasets. In *COLM*, 2024.

356 T. Marty, E. Elmoznino, L. Gagnon, T. Kasetty, M. Nishikawa-Toomey, S. Mittal, G. Lajoie, and
357 D. Sridhar. A compression perspective on simplicity bias. *arXiv preprint arXiv:2603.25839*,
358 2026.

359 K. O. McGraw and S. P. Wong. A common language effect size statistic. *Psychological Bulletin*,
360 111(2):361–365, 1992.

361 A. Meurer, C. P. Smith, M. Paprocki, O. Čertík, S. B. Kirpichev, M. Rocklin, A. Kumar, S. Ivanov,
362 J. K. Moore, S. Singh, T. Rathnayake, S. Vig, B. E. Granger, R. P. Muller, F. Bonazzi, H. Gupta,
363 S. Vats, F. Johansson, F. Pedregosa, M. J. Curry, A. R. Terrel, Š. Roučka, A. Saboo, I. Fernando,
364 S. Kulal, R. Cimman, and A. Scopatz. SymPy: symbolic computing in Python. *PeerJ Computer
365 Science*, 3:e103, 2017. doi: 10.7717/peerj-cs.103.

366 C. Mingard, G. Valle-Pérez, J. Skalse, and A. A. Louis. Is SGD a Bayesian sampler? well, almost.
367 *JMLR*, 2021.

368 N. Nanda, L. Chan, T. Lieberum, J. Smith, and J. Steinhardt. Progress measures for grokking via
369 mechanistic interpretability. In *ICLR*, 2023.

370 Z. Pan, S. Wang, and J. Li. Understanding LLM behaviors via compression: Data generation,
371 knowledge acquisition and scaling laws. *arXiv preprint arXiv:2504.09597*, 2025.

372 A. Paszke, S. Gross, F. Massa, A. Lerer, J. Bradbury, G. Chanan, T. Killeen, Z. Lin, N. Gimelshein,
373 L. Antiga, A. Desmaison, A. Köpf, E. Z. Yang, Z. DeVito, M. Raison, A. Tejani, S. Chilamkurthy,
374 B. Steiner, L. Fang, J. Bai, and S. Chintala. PyTorch: An imperative style, high-performance deep
375 learning library. In *NeurIPS*, 2019.

376 G. Penedo, H. Kydlíček, A. Lozhkov, M. Mitchell, T. Wolf, et al. The FineWeb datasets: Decanting
377 the web for the finest text data at scale. *arXiv preprint arXiv:2406.17557*, 2024.

378 A. Radford, J. Wu, R. Child, D. Luan, D. Amodei, and I. Sutskever. Language models are unsuper-
379 vised multitask learners. Technical report, OpenAI, 2019.

380 S. Ravfogel, G. Yehudai, T. Linzen, J. Bruna, and A. Bietti. Emergence of linear truth encodings in
381 language models. In *NeurIPS*, 2025.

382 J. Rissanen. Modeling by shortest data description. *Automatica*, 14(5):465–471, 1978.

383 D. Rolnick, A. Veit, S. Belongie, and N. Shavit. Deep learning is robust to massive label noise.
384 *arXiv preprint arXiv:1705.10694*, 2017.

385 C. E. Shannon. A mathematical theory of communication. *Bell System Technical Journal*, 27(3):
386 379–423, 1948.

387 R. J. Solomonoff. A formal theory of inductive inference. *Information and Control*, 7(1):1–22,
388 1964.

389 G. Valle-Pérez, C. Q. Camargo, and A. A. Louis. Deep learning generalizes because the parameter-
390 function map is biased towards simple functions. In *ICLR*, 2019.

391 J. Wan and L. Mei. Large language models as computable approximations to Solomonoff induction.
392 *arXiv preprint arXiv:2505.15784*, 2025.

393 Wikimedia Foundation. English Wikipedia dump, 2023-11-01. [https://huggingface.co/
394 datasets/wikimedia/wikipedia](https://huggingface.co/datasets/wikimedia/wikipedia), 2023. Subset 20231101.en.

395 J. Xie, K. Zhang, J. Chen, R. Lou, and Y. Su. Adaptive chameleon or stubborn sloth: Revealing the
396 behavior of large language models in knowledge conflicts. In *ICLR*, 2024.

- 397 A. Yang, A. Li, B. Yang, B. Zhang, B. Hui, B. Zheng, B. Yu, C. Gao, C. Huang, C. Lv, et al. Qwen3
398 technical report. *arXiv preprint arXiv:2505.09388*, 2025.
- 399 C. Zhang, S. Bengio, M. Hardt, B. Recht, and O. Vinyals. Understanding deep learning requires
400 rethinking generalization. In *ICLR*, 2017.

401 A Reproducibility

402 All code, data generation scripts, and evaluation scripts are available at [anonymous repository, in-
403 cluded as supplementary material]. Denoising experiments were run on Modal.com T4 GPUs using
404 PyTorch. Standard math and Wikipedia experiments were run on Apple Mac M4 (36GB) using
405 MLX (v0.31.0). Qwen3-0.6B experiments were run on Modal.com A10G GPUs; the ~1B Qwen3-
406 architecture experiments were run on Modal.com A100 GPUs. Over 210 models were trained across
407 all conditions; total compute approximately 150 GPU-hours.

408 **Corpus generation examples.** Each problem is generated programmatically and verified by
409 SymPy. Below is a representative training example (denoising, equal random condition):

410 Problem: Expand and simplify $3 * (2x + 5) - 4x$
411 Step 1: $3 * (2x + 5) = 6x + 15$
412 Step 2: $6x + 15 - 4x = 2x + 15$
413 Answer: $2x + 15$

414 The corresponding random-error version replaces a derivation step with a plausible but incorrect
415 computation (e.g., $6x + 15 - 4x = 2x + 11$). In the coherent condition, a systematic rule is
416 applied consistently across all problems of the same type (e.g., for distribution: $a(b + c) = ab + c$
417 instead of $ab + ac$). The paired test set presents the same prompt with both correct and incorrect
418 completions for NLL comparison.

419 B Formal MDL Predictions

420 We state the theoretical predictions that motivate our experiments. All three describe the behav-
421 ior of an idealized two-part MDL learner [Rissanen, 1978, Grünwald, 2007], not formal theorems
422 about neural network training dynamics. Our experiments (Sections 4.1–4.3) test whether gradient-
423 descent-trained transformers approximate this behavior.

424 **Setup.** Consider a corpus D containing n problems. In the denoising setup, each problem i appears
425 with two answers: one generated by a correct rule system T_1 and one by an alternative system T_2 .
426 The model must encode *both* answers—it cannot simply “select” one theory and discard the other.
427 We model this as a latent-source mixture: for each problem, a latent variable $z_i \in \{1, 2\}$ indicates
428 which source generated the answer. The model encodes: (1) the mixture structure (source theories +
429 mixing weights), and (2) the data given the mixture, including the per-problem source assignments.

430 The total description length is:

$$L(M) + L(D|M) = \underbrace{L(T_1) + L(T_2)}_{\text{theories}} + \underbrace{L(\{z_i\})}_{\text{source assignments}} + \underbrace{\sum_i L(a_i|T_{z_i})}_{\text{answers given sources}}$$

431 The key quantity is the cost of encoding the alternative system T_2 plus the residuals that T_2 cannot
432 predict. A model that internalizes the more compressible theory can encode answers from that theory
433 cheaply and must pay a higher per-token cost for answers from the less compressible theory—which
434 is exactly what paired evaluation measures (lower NLL on the more compressible cluster).

435 **Prediction 1** (random errors). *If T_1 is a compact rule system with $L(T_1) = O(1)$ and T_2 consists
436 of independent random answers, then correct-answer completions receive lower NLL than incorrect
437 ones.*

438 *Proof sketch.* The correct answers are generated by a compact T_1 : once internalized ($O(1)$ bits),
439 each correct answer costs $O(c)$ bits to encode. The random false answers have no shared structure:
440 encoding each requires memorizing an independent value, contributing $O(n)$ total bits to $L(T_2)$.

The model’s optimal strategy allocates more probability mass to the predictable (correct) cluster, resulting in lower NLL for correct completions.

Prediction 2 (coherent errors). *If both T_1 and T_2 are compact rule systems with $L(T_1) \approx L(T_2) = O(1)$, then at equal frequency the model assigns comparable NLL to both clusters.*

Proof sketch. Both theories have comparable model complexity: $L(T_1) \approx L(T_2)$. Given either theory, the per-answer encoding cost is the same: $L(a_i|T_1) \approx L(a_i|T_2)$ for their respective answers. Since both clusters are equally compressible and equally frequent, the optimal mixture assigns equal probability mass to each, yielding comparable NLL.

Prediction 3 (multi-rule errors). *If the false system consists of N compact rules $\{r_1, \dots, r_N\}$ with random assignment to problems, then encoding the false cluster requires an additional $O(n \log N)$ bits for source selection, restoring the NLL advantage of the correct cluster.*

Proof sketch. Each rule r_i is compact: $L(r_i) = O(1)$. However, encoding which rule applies to which problem requires a selector $s : \{1, \dots, n\} \rightarrow \{1, \dots, N\}$. When assignments are random (i.i.d. uniform), the selector has entropy $n \log_2 N$ bits and is incompressible. The total false-cluster code length is $L(\{r_i\}) + L(s) + \sum_i L(a_i|r_{s(i)}) = O(N) + O(n \log N) + O(n \cdot c)$. For $N \geq 2$, the $n \log N$ term makes the false cluster more expensive to encode than the correct cluster, which requires no selector. Note that this prediction concerns the *total* encoding cost of the false cluster, not only the per-completion probability mass split across N rules (see Section 5 for discussion of this alternative explanation).

C Compression Measure (gzip)

To operationalize the MDL argument, we measure compression ratios (gzip level 9) on concatenated correct vs incorrect completions from each paired test set.

Condition	Correct	Incorrect	Delta	Paired Accuracy
random 50/50	0.1627	0.1639	+0.0012	79.5%
coherent 50/50	0.1656	0.1658	+0.0002	47.2%
contradictory	0.1745	0.1744	−0.0001	49.0%
multirule $N=2$	0.1671	0.1710	+0.0038	77.6%
multirule $N=3$	0.1669	0.1722	+0.0053	82.8%
multirule $N=5$	0.1672	0.1730	+0.0057	84.8%
multirule $N=10$	0.1676	0.1752	+0.0076	88.3%
world random	0.0370	0.0516	+0.0146	57.7%
world coherent	0.0374	0.0384	+0.0010	46.6%

Table 9: Compression ratio (gzip) and paired accuracy across all 9 conditions. Spearman $\rho = 0.68$, $p = 0.042$.

Conditions with larger compression gaps produce stronger truth bias. We present this as supporting evidence; the primary argument rests on the experimental contrasts in Section 4.

D Standard (Non-Denoising) Experiments

D.1 Standard Corpus Design

In the standard setup, each problem appears once with either a correct or incorrect solution. The two groups do not share prompts.

D.2 Standard vs Denoising

The denoising setup is harder (within-problem contradiction vs across-corpus mixing), but the gap closes with scale.

Size	Standard	Denoising	Gap
tiny	79.5%	65.3%	−14.2 pp
small	86.2%	74.6%	−11.6 pp
medium	87.1%	81.1%	−6.0 pp
large	88.3%	85.2%	−3.1 pp

Table 10: Standard vs denoising accuracy (50/50 random). The gap closes with scale.

Proportion	Random	Coherent	False pref.
50/50	80%	47.2%	52.8%
40/60	79%	27.8%	72.2%
30/70	75%	14.7%	85.3%
20/80	69%	9.6%	90.4%
10/90	67%	—	—

Table 11: Random vs coherent accuracy across proportions (standard, tiny). Random: truth bias persists at 10/90. Coherent: the model follows pure frequency. “False pref.” = $1 - \text{coherent accuracy}$, i.e., the fraction of pairs where the model prefers the coherent false completion.

472 D.3 Frequency Effects

473 D.4 Noise Tolerance

Condition	Ratio	Tiny	Small	Medium	Large
Equal random	1:1	65.3%	74.6%	81.1%	85.2%
2:1 noise	1:2	59.0%	68.6%	73.6%	75.2%
4:1 noise	1:4	56.6%	64.9%	66.6%	65.8%

Table 12: Paired accuracy at increasing noise ratios (denoising). Accuracy degrades gracefully; at 4:1 noise, it plateaus at medium/large.

474 D.5 NLL Distribution

475 Random errors produce a right-skewed per-pair NLL difference (mean +0.048, median +0.025,
476 81.5% positive). Coherent errors produce a symmetric distribution (45.5% positive). This explains
477 why pair accuracy and mean ΔLoss can diverge.

478 E Additional Transfer Experiments

479 E.1 Synthetic World

480 A synthetic world with 50 entities, 4 types, 15 deterministic rules. Random 50/50: $57.7\% \pm 1.7\%$.
481 Coherent 50/50: $46.6\% \pm 1.7\%$. Per-type: minerals best (68.7%), potions near chance (49.1%).

482 E.2 Wikipedia Per-Entity-Type

483 E.3 Cross-Domain Falsification

484 Adding correct cross-domain tasks to a coherent corpus selectively increases derivative accuracy
485 ($35\% \rightarrow 56\%$ at 25%). The effect is non-monotonic (drops at 50% due to dilution).

Entity Type	N Pairs	Accuracy
LOC	45	82.2%
NORP	197	78.7%
GPE	301	77.1%
PERSON	402	69.9%
ORG	664	65.2%
CARDINAL	172	61.0%

Table 13: Wikipedia accuracy by entity type (random substitution, averaged across sizes).

Setup	Tokenizer	Random	Coherent
Standard	Char	79.5%	47.2%
Standard	BPE	85.6%	45.9%
Denoising	Char	65.3%	43.5%
Denoising	BPE	75.9%	49.3%

Table 14: BPE vs char-level (tiny, 4 seeds). The effect survives BPE and strengthens for random errors.

F Robustness Checks

F.1 BPE Tokenization

F.2 Chained Verification

Embedding verification dependencies in coherent-error tasks restores truth bias: 70.9% at tiny (vs 43% standard coherent). The effect *decreases* with model size (64% at small, 61% at large)—larger models absorb coherent errors more readily, making verification less effective.

F.3 Learning Curves

All sizes reach behavioral plateau by step 3000–4000. The large model achieves 88.8% at step 3000, stable through 5000.

G Per-Seed Details

Tables 15–16 report per-seed accuracy for the core denoising conditions. Per-seed breakdowns for multi-rule, Wikipedia, and Qwen3 experiments are available in the supplementary code repository (evaluation logs per model).

Size	Seed 1	Seed 2	Seed 3	Seed 4	Mean
tiny	64.2%	64.1%	67.0%	65.8%	65.3%
small	75.1%	76.3%	72.6%	74.2%	74.6%
medium	79.6%	80.9%	82.4%	81.3%	81.1%
large	83.5%	86.8%	–	–	85.2%

Table 15: Denoising equal random, per-seed accuracy.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [\[Yes\]](#)

Size	Seed 1	Seed 2	Seed 3	Seed 4	Mean
tiny	40.2%	44.3%	46.3%	43.0%	43.5%
small	44.8%	45.3%	40.3%	47.4%	44.5%
medium	45.8%	40.9%	47.7%	48.6%	45.8%
large	51.6%	50.4%	–	–	51.0%

Table 16: Denoising equal coherent, per-seed accuracy.

Justification: All claims are explicitly qualified to our controlled experimental setting. The abstract states the scope (3.5M–1B parameters, controlled corpora) and the introduction notes that extension to large-scale pretraining is an open question.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: A dedicated Limitations section discusses model scale, domain specificity, discriminative vs generative gap, seed counts, and causal identification.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The paper does not claim formal theoretical results. Appendix B presents three MDL-motivated predictions with proof sketches for an idealized learner, explicitly marked as informal predictions rather than formal theorems about neural network training. The paper is primarily empirical.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: All hyperparameters, model architectures, data generation procedures, evaluation protocols, and random seeds are specified in Section 3 and Appendix A. Full generation templates are included in supplementary material.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).

- (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: Anonymous repository included as supplementary material with full code for data generation, training, and evaluation.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: See Section 3.2 (optimizer, lr, schedule, batch size, steps) and Appendix A (corpus examples, hardware). Train/test separation is guaranteed by seed: test problems use a different seed (999) than training.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: All $\pm$ values are standard deviations across random seeds (not standard errors). p -values are computed via Wilcoxon signed-rank test (one-sided) on per-pair NLL differences; bootstrap 95% CIs (10,000 resamples) are used for confidence intervals. The

number of seeds is reported for every condition. Per-seed breakdowns for the core denoising conditions are provided in Appendix G. For the 1B experiment (single seed), the p -value reflects per-example significance and is marked as descriptive.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: See Appendix A: approximately 150 GPU-hours total on Modal.com T4/A10G/A100 GPUs and Apple M4 Pro (36GB). Hardware is specified per experiment type.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The research uses only synthetic data and publicly available Wikipedia text. No human subjects, no private data, no dual-use concerns beyond those discussed in Ethical Considerations.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.

- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: The Ethical Considerations section discusses both the defensive value (informing data curation) and the risk (characterizing conditions under which misinformation evades filtering).

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: All models are small ($\leq 1B$ parameters) and trained on synthetic or filtered web data. They have no capabilities beyond the narrow experimental domain and pose no misuse risk.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: Wikipedia data is used under its CC BY-SA license. FineWeb-Edu [Penedo et al., 2024] is used under its ODC-By license. All software dependencies (PyTorch, MLX, SymPy, spaCy) are open source.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: Code and data generation scripts are documented in the supplementary material with README and inline comments.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: No human subjects or crowdsourcing involved. All data is generated programmatically.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.

812 • According to the NeurIPS Code of Ethics, workers involved in data collection, cura-
813 tion, or other labor should be paid at least the minimum wage in the country of the
814 data collector.

815 **15. Institutional review board (IRB) approvals or equivalent for research with human**
816 **subjects**

817 Question: Does the paper describe potential risks incurred by study participants, whether
818 such risks were disclosed to the subjects, and whether Institutional Review Board (IRB)
819 approvals (or an equivalent approval/review based on the requirements of your country or
820 institution) were obtained?

821 Answer: [N/A]

822 Justification: No human subjects involved.

823 Guidelines:

824 • The answer [N/A] means that the paper does not involve crowdsourcing nor research
825 with human subjects.

826 • Depending on the country in which research is conducted, IRB approval (or equiva-
827 lent) may be required for any human subjects research. If you obtained IRB approval,
828 you should clearly state this in the paper.

829 • We recognize that the procedures for this may vary significantly between institutions
830 and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the
831 guidelines for their institution.

832 • For initial submissions, do not include any information that would break anonymity
833 (if applicable), such as the institution conducting the review.

834 **16. Declaration of LLM usage**

835 Question: Does the paper describe the usage of LLMs if it is an important, original, or
836 non-standard component of the core methods in this research? Note that if the LLM is used
837 only for writing, editing, or formatting purposes and does *not* impact the core methodology,
838 scientific rigor, or originality of the research, declaration is not required.

839 Answer: [N/A]

840 Justification: LLMs (Claude) were used for writing assistance, code editing, and bibliog-
841 raphy verification, but not as a component of the research methodology. All experimental
842 code, data generation, model training, and evaluation were implemented and verified inde-
843 pendently of LLM assistance.

844 Guidelines:

845 • The answer [N/A] means that the core method development in this research does not
846 involve LLMs as any important, original, or non-standard components.

847 • Please refer to our LLM policy in the NeurIPS handbook for what should or should
848 not be described.